

Automatic Detection of Blood Vessels From Retinal Images Using Convolutional Neural Network

Ville Virkkala
(Dated: May 15, 2019)

Main causes of blindness, such as diabetic retinopathy and hypertension, are related to morphological changes of the blood vessels. Thus automatic segmentation of blood vessels from retinal images could greatly help to make the correct diagnosis efficiently in optometric studies. In this work a convolutional neural network is developed to automatically segment blood vessels from retinal images. The developed method produces an accuracy of 85% for the test set. The performance of developed method is also compared against other similar methods and against random forest classifier.

I. INTRODUCTION

Diabetic retinopathy is the leading cause of vision loss among working age adults. At the first stage of diabetic retinopathy there is no symptoms and the only way to diagnose the disease is by fundus photography. Diabetic retinopathy affects the structure of retinal vasculature causing identifiable attributes in retinal blood vessels such as width, tortuosity and branching pattern¹. Untreated the diabetic retinopathy can lead to nearly complete blindness where patient can only tell the light from darkness. In addition to diabetic retinopathy other eye diseases, like hypertension caused by the damage to the retina and retinal circulation by high blood pressure, can be diagnosed by examining the retinal blood vessels. One in eleven adults, 435 million, are diagnosed with diabetes² and thus there is a huge demand for automatic retinal vessel segmentation to save medical resources.

Several studies have used neural networks in automatic blood vessel detection from retinal images. These methods can be roughly divided into two groups: patch based segmentation of blood vessels³ and methods based on fully convolutional neural networks^{4,5}. In patch based image segmentation the image is traversed through pixel by pixel. For each pixel a patch of fixed size, centered at the pixel, is taken from the image and fed to neural network that classifies the pixel into certain class. In Ref.³, for each pixel three patches of different sizes from the green channel of image were taken and scaled into same size and used as a three channel input to convolutional neural network. In Ref.³ the reported accuracy was in most cases clearly better than 90%. The advantages of patch based method are their simplicity and ease of training. However, their main disadvantage is the high computational load when doing inference, because for each pixel separate patch is taken that is fed to neural network. Fully convolutional neural networks⁴ are the current state of the art method in image segmentation. The huge advantage of fully convolutional neural network is the huge speed up compared to patch based methods, because the whole image is fed only once as whole to the neural network. In addition there is more contextual information available in fully convolutional neural networks because the whole image is processed at once instead of

using smaller patches. In Ref.⁵ a modified version of the fully convolutional network in Ref.⁴ was used. In Ref.⁵ the downsampling of the image was done using the same VGG-16 network as in⁴ however the output of each convolution stage was directly upsampled into original size and a binary cross entropy loss was then connected into each upsampled image. The final loss that was optimized is then the sum of the losses connected to each upsampled image. In Ref.⁵ the reported accuracy was in all cases better than 95% while the method was much faster, single image segmentation time around 11seconds, than the method of Ref.³, single image segmentation time around 2000 seconds.

In this work a patch based semantic segmentation method based on convolutional neural networks, resembling that of Ref.³, is developed to detect blood vessels from retinal images. The performance of the developed method is validated against several test images. In addition performance of the developed method is compared to random forest classifier implemented in scikit-learn python library⁶. The paper is organized as follows. The used data-set and the computational methods are described in detail in sections II and III. In Sec. IV training metrics and performance of the developed method in inference when applied on test images are given. Sec. V is a summary of the results and the differences between the developed method and the reference random forest classifier are discussed.

II. USED DATA-SET

The used data set in this work is the publicly available DRIVE data set⁷ available in⁸. The data set consist of 40 images of size 565x584 and the corresponding ground truth annotations of blood vessels. The images were obtained from diabetic subjects as a part of diabetic retinopathy screening program in The Netherlands.

The data set was divided into training, validation and test sets containing 28, 6 and 6 images respectively. The images are three channel RGB images, but only the green channel is used in classification, because of the high contrast between blood vessels and background in green channel⁹. In classification of pixels patches of size 33x33

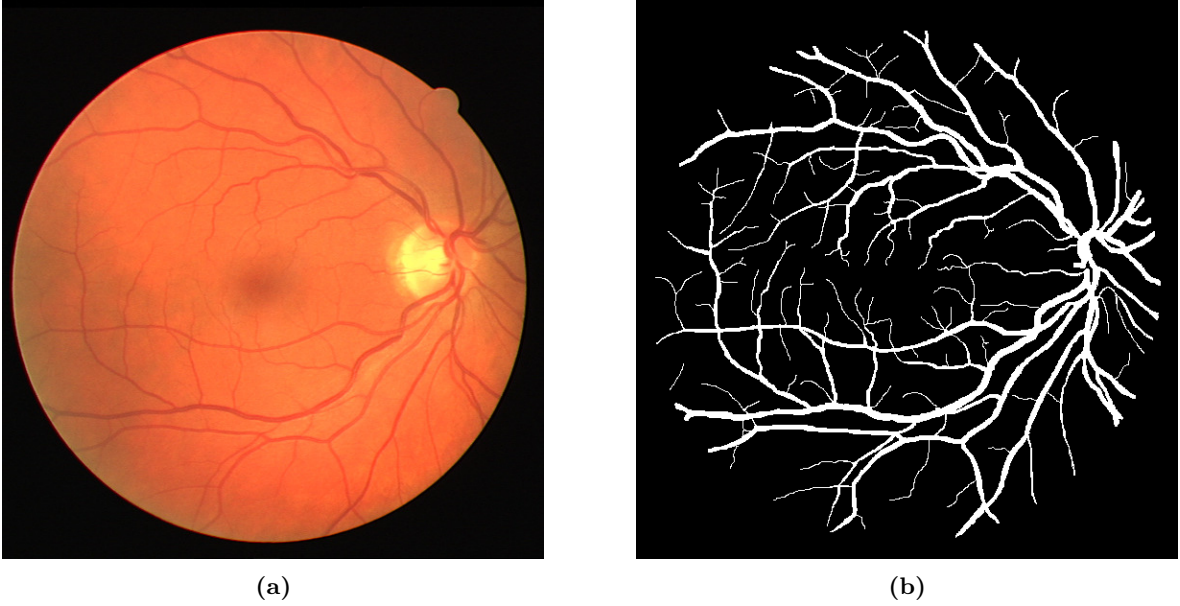


FIG. 1: Example of retinal image used in training (a) and the corresponding ground truth annotation of blood vessels (b).

centered at the pixel were used. From each image 2000 patches for both pixels corresponding to blood vessels and background pixels were randomly sampled resulting in total of 112000 samples for training set, 24000 samples for validation set and 24000 samples for test set respectively. Examples of retinal image that is classified and the corresponding ground truth image are shown in figures 1a and 1b respectively. Examples of patches that corresponds to pixels that are labeled as blood vessels and background are shown in figures 2a-2c and 2d-2f respectively.

III. COMPUTATIONAL METHODS

The developed neural network in this work is implemented using the Python programming language and the PyTorch¹⁰ open-source machine learning library for Python. The reference, random-forest classifier, method is implemented using the Scikit-learn machine learning library⁶. The source code for most relevant parts of the algorithm are given in Appendix A

A. Developed Convolutional Neural Network

The image is processed pixel by pixel and for each pixel a patch of size 33x33, centered at the pixel, is taken from the image and used in classification. The developed convolutional neural network resembles that of Ref.³. However, instead of using three different patch resolutions in classification for a given pixel, only one resolution corresponding to 33x33 pixels from green channel is used in this work. The architecture of the developed neural net-

work is shown in figure 3. The input of size 33x33 pixels is fed to two convolutional layers with kernels of sizes 5x5. Each convolutional layer is followed by linear rectifying unit (RELU), a patch normalization layer and a max-pool layer with stride 2. The output of second max-pool layer is flattened and fed to fully connected layer followed by RELU. The RELU is followed by dropout layer with dropout probability of 0.5. Finally the output of the dropout layer is forwarded to last fully connected layer with single output. The output of the network is converted to probability using sigmoid function and the final class label is obtained using round operator.

B. Neural Network Training

The developed method above contains three hyper parameters, the number of kernels in first and second convolution layer and the number of layers in first fully connected layer. Rest of the parameters, *i.e.*, the kernel size and patch size were kept fixed. The optimal hyper parameters were found using random search. Ten different configurations were randomly sampled from given intervals that were [10, 35], [10, 35] and [200, 1200] respectively for the hyper parameters. For each configuration the network was trained over 125 epochs using the train data set. In addition an early stopping criterion with tolerance of 0.01 and patience of 2 was used to avoid over fitting of the model. At the end of each epoch the validation loss, evaluated against the validation data set, was evaluated and used as input for early stopping criterion.

Stochastic gradient descent with patch size of 1000 and Adam optimizer¹¹, initialized with learning rate of 1e-6 was used to optimize the network parameters. The final accuracy of the network, with given hyper param-

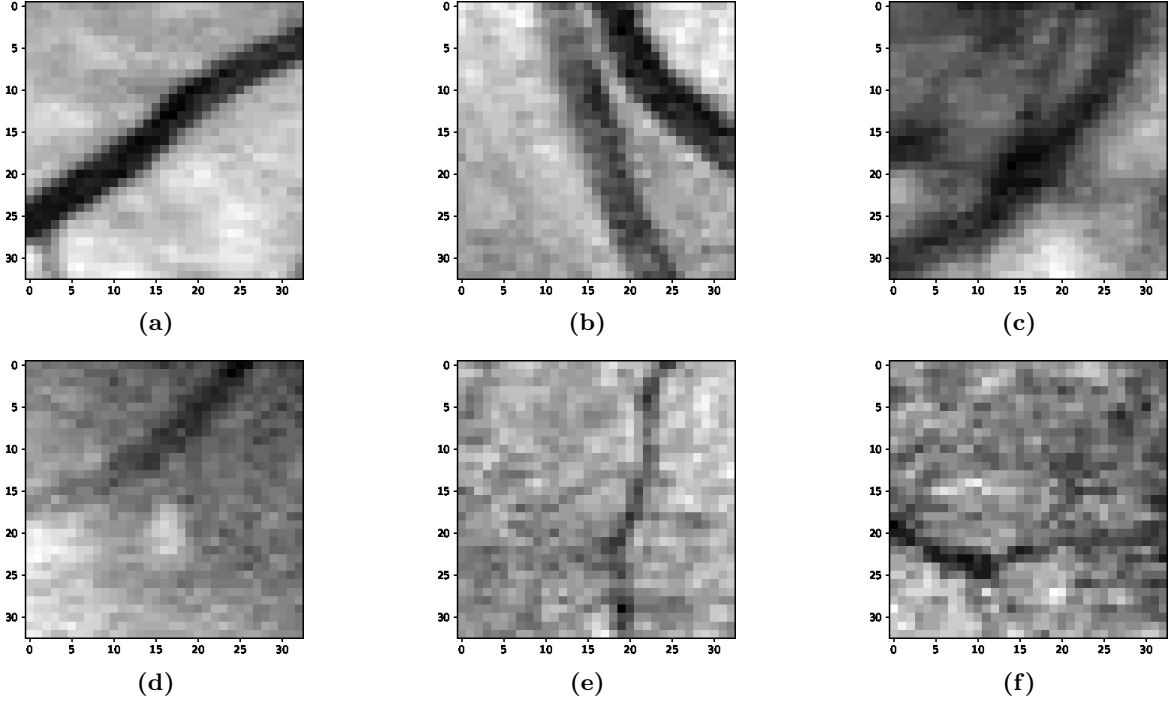


FIG. 2: Examples of patches corresponding to pixels that are labeled as blood vessels (a)-(c) and patches corresponding to pixels that are labeled as background (d)-(f). The pixel that the label is related to is located at the center of patch.

ters, was then obtained evaluating the accuracy against the test set. Final network, used in inference, was then constructed using the hyper parameters corresponding to best accuracy obtained. The final network was then re-trained over 250 epochs or until early stop to get full convergence.

C. Random forest classifier training

The random forest classifier, that was used as a reference method, was implemented and trained using the Scikit-learn machine learning Python library. In used random forest classifier 200 estimators were used. The classifier was trained using the whole training data set.

D. Inference

When evaluating the performance of the developed neural network, inference is performed on each image in test data set. After inference accuracy, sensitivity and specificity metrics are evaluated. The metrics are evaluated by comparing the combined predicted labels, including all six images in test data set, to the true labels indicated by the ground truth annotations. The accuracy, sensitivity and specificity are calculated as below

$$SENSITIVITY = \frac{TP}{TP + FN}, \quad (1)$$

$$SPECIFICITY = \frac{TN}{FP + TN}, \quad (2)$$

$$ACCURACY = \frac{TP + TN}{TP + FP + TN + FN}, \quad (3)$$

where TP is the number of true positives, FN is the number of false negatives, TN is the number of true negatives and FP is the number of false positives. The performance metrics are evaluated for both developed neural network and for the trained random forest classifier that is used as a reference method.

IV. RESULTS

A. Optimal hyper parameters and the training accuracy for developed neural network

The results for the hyper parameter search, explained in section III B, are shown in Table I. According to Table I the best accuracy was obtained when the hyper parameters were set to 33x67x355. The final accuracy with these parameters for test set, after training over 250 epochs was 84.8%.




FIG. 3: Architecture of the developed convolutional neural network.

B. Training accuracy of the random forest classifier

The obtained accuracy for the test data set using the trained random-forest classifier was around 88%. Thus the obtained accuracy for the random-forest classifier is somewhat better than for the developed neural network.

C. Inference

The inference was performed on each image in test data set and combined performance metrics were evaluated as described in III D. The summary of the results for both developed method and the reference method are shown in table. For developed method it took around 1 minute to perform inference for a single image while for the reference method the corresponding time was around 30s being half of the computational time of the developed method. In figures 4(a)-4(d) and 4(e)-4(h) results of inference in the case of worst and best accuracy are shown respectively for both developed neural network and for

reference random forest classifier. The first image in the row is the retinal image belonging to test set, second image is the manual ground truth of blood vessels, third image contains the predicted blood vessels using the developed neural network and the last image contains the predicted blood vessels using the reference random forest classifier.

Method	Acc. (%)	Sens. (%)	Spec. (%)
Neural Network	91	85	91
Random Forest	92	83	93

TABLE II: Performance metrics, accuracy, sensitivity and specificity, for developed neural network and reference random forest classifier.

Configuration	Accuracy (%)	Configuration	Accuracy (%)
33x67x355	83.8	16x50x981	82.9
29x42x709	83.7	18x63x1048	82.9
29x49x871	83.5	27x58x359	82.4
19x57x1158	83.3	10x64x1176	81.3
20x60x256	83.2	14x49x638	80.1

TABLE I: Results of the random hyper parameter search.

The numbers in the configuration column indicates the number of kernels in the first and second convolution layers and the number of neurons in the first fully connected layer respectively. The accuracy is obtained running the trained network with the given hyper parameters against the test data set.

For developed neural network the accuracy for test images ranged from 89% to 93% the combined accuracy being 91%. The confusion matrix corresponding to combined results for neural network classifier is shown in figure 5. Based on the confusion matrix the calculated sensitivity and specificity, equations 1 and 2, for the developed neural network classifier are 85.3% and 91.4% respectively.

For reference random forest classifier the accuracy for test images ranged from 89% to 94% the combined accuracy being 92%, which is slightly better than for developed method. The confusion matrix corresponding to combined results for reference random forest classifier is shown in figure 6. Based on the confusion matrix the sensitivity and specificity for the reference method are 82.9% and 92.5% respectively.

According to figure 4 both methods works relatively well capturing the essential characteristics of blood vessels. However for developed neural network somewhat




FIG. 4: Results of inference in the case of worst accuracy (a)-(d), (a) the retinal image, (b) the manual annotation, (c) result of inference using the develop neural network and (d) the result of inference when the random forest classifier is used. (e)-(h) results of inference in the case of best accuracy.

more noise is visible in predicted blood vessels compared to reference method which is also evident from the evaluated specificities. In addition the developed method has some problems in the boundary region separating the actual image from the background where multiple false positives are visible. At first glance it is surprising that the evaluated accuracy is clearly better for inference compared to training. However this can be explained by the fact that in training the positive and negative samples were perfectly balanced. However in inference the whole image, containing significantly more negatives samples,

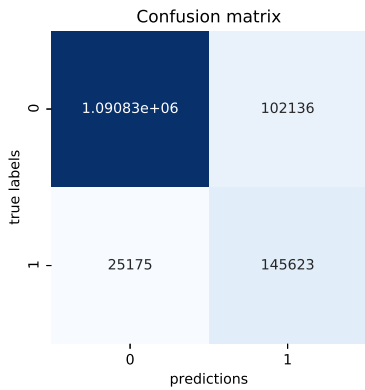


FIG. 5: Confusion matrix corresponding to inference results for developed neural network.




FIG. 6: Confusion matrix corresponding to inference results for trained random forest classifier.

was classified. Both method can classify more accurately negative samples, which is evident from the evaluated sensitivities and specificities, and thus the metrics show better performance in inference compared to training.

V. CONCLUSIONS

In this work a convolutional neural network was developed to automatically detect blood vessels from retinal images. The obtained accuracy, sensitivity and speci-

ficity for the developed method were 91%, 85% and 91% respectively. In addition the a reference performance metrics were evaluated using a random forest classifier for which the corresponding metrics were 92%, 83% and 93%. It is good to notice that the performance metrics are somewhat affected by the class imbalance, favoring negative samples, in inference leading to overestimated performance. The performance of the developed method is somewhat worse compared to reference method even the computational speed of the reference method is about the double of the developed method. However the sensitivity of the developed method is still better than for the reference method.

In Ref.³ a rather similar neural network was developed, the main difference being that three different patch sized were used in classification instead of just one. In Ref.³

they reported somewhat better performance compared to developed method however the reported computational time in Ref.³ was nearly 60 minutes for single image, whereas for developed method the computational time was only 1 minute per image. It is not clear what explains the huge difference in computational time between the two method because in both methods the performance was evaluated using cpu.

Even the evaluated performance for the developed method is slightly worse compared to reference method the neural networks can still be considered as the better approach in a long run. The neural networks can be easily extended and for deep networks, like fully convolutional networks used in Ref.⁵, significant improvements in both performance and computational time can be achieved that outperforms other classifiers.

¹ C. Zhu, B. Zou, R. Zhao, J. Cui, X. Duan, Z. Chen, and Y. Liang, Computerized Medical Imaging and Graphics **55**, 68 (2017), special Issue on Ophthalmic Medical Image Analysis.

² <https://diabetesatlas.org/>.

³ J. H. Tan, U. R. Acharya, S. V. Bhandary, K. C. Chua, and S. Sivaprasa, Journal of Computational Science **309**, 70 (2018).

⁴ J. Long, E. Shelhamer, and T. Darrell, in *2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* (2015) pp. 3431–3440.

⁵ K. Hu, Z. Zhang, X. Niu, Y. Zhang, C. Cao, F. Xiao, and X. Gao, Neurocomputings **105**, 179 (2009).

⁶ F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cour-

napeau, M. Brucher, M. Perrot, and E. Duchesnay, Journal of Machine Learning Research **12**, 2825 (2011).

⁷ J. Staal, M. Abramoff, M. Niemeijer, M. Viergever, and B. van Ginneken, IEEE Transactions on Medical Imaging **23**, 501 (2004).

⁸ <http://https://www.isi.uu.nl/Research/Databases/DRIVE//>.

⁹ A. Elbalaoui, M. Fakir, K. Taifi, and A. Merbouha, in *2016 13th International Conference on Computer Graphics, Imaging and Visualization (CGiV)* (2016) pp. 324–332.

¹⁰ A. Paszke, S. Gross, S. Chintala, G. Chanan, E. Yang, Z. DeVito, Z. Lin, A. Desmaison, L. Antiga, and A. Lerer, in *NIPS-W* (2017).

¹¹ D. P. Kingma and J. Lei Ba, (2015), arXiv:1412.6980.

Appendix A: Source Code

1. Convolutional Network

```

import torch.nn as nn
import torch.nn.functional as F

class ConvNet(nn.Module):
    def __init__(self, in_channels, out_channels1, out_channels2, out_dens):
        super(ConvNet, self).__init__()
        self.out_channels2 = out_channels2
        self.conv1 = nn.Conv2d(in_channels, out_channels1, kernel_size=5,
                                stride=1, padding=0)
        self.batchn1 = nn.BatchNorm2d(out_channels1)
        self.pool1 = nn.MaxPool2d(kernel_size=2, stride=2, padding=0)
        self.conv2 = nn.Conv2d(out_channels1, out_channels2, kernel_size=5,
                                stride=1, padding=0)
        self.batchn2 = nn.BatchNorm2d(out_channels2)
        self.pool2 = nn.MaxPool2d(kernel_size=2, stride=2, padding=0)

        self.fc1 = nn.Linear(out_channels2 * 5 * 5, out_dens)
        self.drop1 = nn.Dropout()
        self.fc2 = nn.Linear(out_dens, 1)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = self.batchn1(x)
        x = self.pool1(x)
        x = F.relu(self.conv2(x))
        x = self.batchn2(x)
        x = self.pool2(x)
        x = x.view(-1, self.out_channels2 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = self.drop1(x)
        x = self.fc2(x)
        return self.sigmoid(x)

```

2. Training

```

from os.path import join
import time
import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim
import sys
import argparse

from conv_models.convnet import ConvNet
from utils.loading import load_samples
from utils.regularization import EarlyStopping
from utils.loss_utils import compute_loss, compute_accuracy
from utils.parameter_search import random_search

```

```

def train(num_epochs, learning_rate, model_params, batch_size, x_train,
          y_train, x_valid, y_valid, x_test, y_test,
          device='cpu', save_model=False):

    net = ConvNet(1, model_params[0], model_params[1], model_params[2])
    net.to(device)

    criterion = nn.BCELoss()

    optimizer = optim.Adam(net.parameters(), lr=learning_rate)
    train_losses = []
    val_errors = []
    val_losses = []

    early_stop = EarlyStopping(tolerance=0.01, patience=2)

    for epoch in range(num_epochs):
        start_time = time.time()
        epoch_loss = 0
        for k in range(int(y_train.size / batch_size) - 1):
            start_ind = k * batch_size
            end_ind = (k + 1) * batch_size if (k + 1) * batch_size <
                y_train.size else y_train.size
            x = torch.tensor(x_train[start_ind:end_ind, :], device=device,
                             dtype=torch.float)
            y = torch.tensor(y_train[start_ind:end_ind, :], device=device,
                             dtype=torch.float)

            # zero the parameter gradients
            optimizer.zero_grad()

            # forward + backward + optimize
            outputs = net(x)
            loss = criterion(outputs, y)
            epoch_loss += np.asscalar(loss.cpu().data.numpy())

            loss.backward()
            optimizer.step()

        # Print accuracy after every epoch
        train_losses.append(epoch_loss)
        validation_accuracy = compute_accuracy(net, x_valid, y_valid)
        val_errors.append(validation_accuracy)
        validation_loss = compute_loss(net, x_valid, y_valid)
        val_losses.append(validation_loss)
        time_taken = (time.time() - start_time)
        print('Accuracy of the network on epoch %d%%' % epoch + ' : %f%%' %
              (100 * validation_accuracy) +
              'validation loss: %f' % validation_loss + 'train loss: %f' %
              epoch_loss + 'took %f' % time_taken + 'seconds')

        if early_stop.stop_criterion(val_losses):
            print('Stop after %d epochs' % epoch)
            break

    test_accuracy = compute_accuracy(net, x_test, y_test)
    if save_model:
        save_filename = join('saved_models',

```



```

'convnet_' + str(model_params[0]) + 'x' + str(model_params[1]) + 'x' +
str(model_params[2]) + '.pth')
torch.save(net.state_dict(), save_filename)
np.savez(join('saved_models', 'training_losses.npz'),
train_losses=train_losses,
val_losses=val_losses,
final_accuracy=test_accuracy)

print('Final_test_accuracy:_%f_%%' % (100 * test_accuracy))
return test_accuracy

```

```

def optimize_hyper_parameters(num_epochs, learning_rate,
                              num_combinations, batch_size,
                              x_train, y_train, x_valid,
                              y_valid, x_test, y_test):

    range1 = [10, 35]
    range2 = [35, 70]
    range3 = [200, 1200]

    parameter_combinations = random_search(num_combinations, range1, range2, range3)
    hyper_parameters = []
    accuracies = []

    for n1, n2, n3 in parameter_combinations:
        accuracy = train(num_epochs, learning_rate, [n1, n2, n3], batch_size,
                        x_train, y_train, x_valid, y_valid, x_test, y_test)
        hyper_parameters.append([n1, n2, n3])
        accuracies.append(accuracy)
        print('accuracy_with_parameters_' + str(n1) + ',_' + str(n2) +
            ',_' + str(n3) + ']' + str(accuracy))

    hyper_parameters = np.array(hyper_parameters)
    accuracies = np.array(accuracies)
    np.savez(join('saved_models', 'hyperparameters4.npz'),
    hyperparameters=hyper_parameters,
    accuracies=accuracies)

    ix = accuracies.argsort()[−1::−1]

    print(accuracies[ix])
    print(hyper_parameters[ix, :])

```

```

def train_random_forests_classifier(x_train, y_train, x_test, y_test, save_model=False):

    from sklearn.ensemble import RandomForestClassifier
    from sklearn.metrics import accuracy_score
    import pickle

    x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[2] *
                                    x_train.shape[3]))
    x_test = np.reshape(x_test, (x_test.shape[0], x_test.shape[2] *
                                   x_test.shape[3]))
    y_train = np.reshape(y_train, (y_train.size, ))
    y_test = np.reshape(y_test, (y_test.size, ))
    classifier = RandomForestClassifier(n_estimators=200, verbose=1, n_jobs=−1)

```

```

classifier.fit(x_train, y_train)
pred_test = classifier.predict(x_test)
rf_accuracy = accuracy_score(y_test, pred_test)
print("Accuracy_of_random_forest: {:.2f}".format(rf_accuracy))

if save_model:
    pickle.dump(classifier, open('saved_models/random_forest.p', 'wb'))

classifier = pickle.load(open('saved_models/random_forest.p', 'rb'))
pred_test = classifier.predict(x_test)
rf_accuracy = accuracy_score(y_test, pred_test)
print("Accuracy_of_random_forest: {:.2f}".format(rf_accuracy))

def main(args):

    parser = argparse.ArgumentParser()
    parser.add_argument('--num_epochs', default=1, type=int)
    parser.add_argument('--batch_size', default=1000, type=int)
    parser.add_argument('--learning_rate', default=1e-6, type=float)
    parser.add_argument('--optimize_hyperparameters', default=False, type=bool)
    parser.add_argument('--save_model', default=False, type=bool)
    parser.add_argument('--num_combinations', default=10, type=int)
    parser.add_argument('--use_random_forest', default=False, type=bool)
    args = parser.parse_args(args)

    use_random_forest = args.use_random_forest
    num_epochs = args.num_epochs
    batch_size = args.batch_size
    learning_rate = args.learning_rate
    hyper_parameter_optimization = False#args.optimize_hyperparameters
    num_combinations = args.num_combinations
    save_model = args.save_model

    (x_train, y_train), (x_valid, y_valid), (x_test, y_test) = load_samples(2000)

    if use_random_forest:
        train_random_forest_classifier(x_train, y_train, x_test, y_test,
                                       save_model)
    elif hyper_parameter_optimization:
        optimize_hyperparameters(num_epochs, learning_rate, num_combinations,
                                  batch_size, x_train, y_train, x_valid, y_valid,
                                  x_test, y_test)
    else:
        train(num_epochs, learning_rate, [33, 67, 355], batch_size, x_train,
              y_train, x_valid, y_valid, x_test, y_test, device='cpu',
              save_model=save_model)

if __name__ == '__main__':
    main(sys.argv[1:])

```

3. Inference

```

from os import listdir, mkdir
from os.path import isfile, isdir, join
import matplotlib.pyplot as plt
import matplotlib.image as mimg

```

```

import numpy as np
import pickle

import torch

from conv_models.convnet import ConvNet
import utils.image_utils as image_utils

def evaluate_performance():

    image_path = join('data', 'DRIVE', 'test', 'images')
    mask_path = join('data', 'DRIVE', 'test', 'mask')
    result_path = join('data', 'results', 'rf')

    images = [f for f in listdir(image_path) if isfile(join(image_path, f))]
    patch_size = 2000

    vessel_detector = VesselDetectorRf(patch_size)

    for image in images:

        id = image[0:2]
        image_file = join(image_path, image)
        mask_file = join(mask_path, id + '_test_mask.gif')
        image_data = mimg.imread(image_file)[: , : , 1]
        mask_data = mimg.imread(mask_file)
        predicted_data = inference(image_data, mask_data, vessel_detector)
        plt.imsave(join(result_path, image), predicted_data, cmap='binary')
        # plt.figure()
        # plt.imshow(image_data, cmap='gray')
        # plt.figure()
        # plt.imshow(predicted_data, cmap='binary', interpolation='nearest')
        # plt.figure()
        # plt.imshow(mask_data, cmap='binary', interpolation='nearest')
        # plt.show()

def inference(image_data, mask_data, vessel_detector):

    indices = np.where(mask_data == 255)
    labels = vessel_detector.detect_vessels(image_data, indices)
    predicted = np.zeros_like(image_data)
    predicted[indices] = labels[:, 0]
    return predicted

class VesselDetector:

    def __init__(self, patch_size, convmodel):
        self.patch_size = patch_size
        self.net = ConvNet(1, 33, 67, 355)
        self.net.load_state_dict(torch.load(convmodel, map_location=lambda storage, loc: storage))

    def detect_vessels(self, image_data, indices):

        self.net.eval()

```

```

labels = np.zeros((indices[0].size, 1), dtype=np.int32)
for k in range(int(np.ceil(indices[0].size / self.patch_size))):
    start_ind = k * self.patch_size
    end_ind = (k + 1) * self.patch_size
    if end_ind > indices[0].size:
        end_ind = indices[0].size

    patch_indices = (indices[0][start_ind:end_ind],
                     indices[1][start_ind:end_ind])
    patch_data = image_utils.get_clipped_area(image_data,
                                              patch_indices)
    patch_data = np.reshape(patch_data, [patch_data.shape[0],
                                         1, patch_data.shape[1],
                                         patch_data.shape[2]])
    x = torch.tensor(patch_data, device='cpu', dtype=torch.float)
    probabilities = self.net(x)
    predicted = torch.round(probabilities.data)
    labels[start_ind:end_ind] = predicted.data.numpy()
    print('Processed_{}_{}%'.format(100 * k / int(np.ceil(indices[0].size
    / self.patch_size))))

return labels

```

```

class VesselDetectorRf:

```

```

    def __init__(self, patch_size):
        self.patch_size = patch_size
        self.classifier = pickle.load(open('saved_models/random_forest.p', 'rb'))

    def detect_vessels(self, image_data, indices):

        labels = np.zeros((indices[0].size, 1), dtype=np.int32)
        for k in range(int(np.ceil(indices[0].size / self.patch_size))):
            start_ind = k * self.patch_size
            end_ind = (k + 1) * self.patch_size
            if end_ind > indices[0].size:
                end_ind = indices[0].size

            patch_indices = (indices[0][start_ind:end_ind],
                             indices[1][start_ind:end_ind])
            patch_data = image_utils.get_clipped_area(image_data,

            patch_data = np.reshape(patch_data, [patch_data.shape[0],
                                                 patch_data.shape[1] *
                                                 patch_data.shape[2]])
            predicted = self.classifier.predict(patch_data)
            labels[start_ind:end_ind, 0] = predicted
            print('Processed_{}_{}%'.format(100 * k / int(np.ceil(indices[0].size
            / self.patch_size))))

        return labels

```

```

if __name__ == '__main__':
    evaluate_performance()

```